

Part_I_exploration_template

November 17, 2022

1 Part I - Dataset Exploration for Ford GoBike System Data

1.1 by Sarah Ali

1.2 Introduction

This data set includes information about individual rides made in a bike-sharing system covering the greater San Francisco Bay area. For this analysis I went ahead to gather datasets for the first year quarter January - April 2019 from this link: <https://s3.amazonaws.com/baywheels-data/index.html>

Ford GoBike, like other bike share systems, consists of a fleet of specially designed, sturdy and durable bikes that are locked into a network of docking stations throughout the city. The bikes can be unlocked from one station and returned to any other station in the system, making them ideal for one-way trips. People use bike share to commute to work or school, run errands, get to appointments or social engagements and more. It's a fun, convenient and affordable way to get around

1.3 Preliminary Wrangling

```
In [1]: # import all packages needed for exploration
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb
```

```
%matplotlib inline
```

```
In [27]: # merge all four csv files into one for the first quarter.
```

```
ford_data = pd.concat(
    map(pd.read_csv, ['January-fordgobike-tripdata.csv', 'Feb-fordgobike-tripdata.csv',
                     'April-fordgobike-tripdata.csv']), ignore_index=True,
    sort=False)
ford_data.sample(5)
```

```
Out[27]:
```

	duration_sec	start_time	end_time
89574	1592	2019-01-20 20:38:21.4610	2019-01-20 21:04:53.8680
192814	1939	2019-02-28 20:08:56.7470	2019-02-28 20:41:16.5700
210165	496	2019-02-26 19:26:11.3040	2019-02-26 19:34:27.6320

```

261132          736  2019-02-19 20:15:46.5840  2019-02-19 20:28:02.7400
690788          210  2019-04-22 09:00:22.7580  2019-04-22 09:03:53.5620

```

```

          start_station_id          start_station_name \
89574          6.0          The Embarcadero at Sansome St
192814         160.0          West Oakland BART Station
210165         64.0          5th St at Brannan St
261132         15.0  San Francisco Ferry Building (Harry Bridges Pl...
690788         44.0  Civic Center/UN Plaza BART Station (Market St ...

```

```

          start_station_latitude  start_station_longitude  end_station_id \
89574          37.804770          -122.403234          17.0
192814          37.805318          -122.294837          274.0
210165          37.776754          -122.399018          60.0
261132          37.795392          -122.394203          66.0
690788          37.781074          -122.411738          58.0

```

```

          end_station_name \
89574  Embarcadero BART Station (Beale St at Market St)
192814          Oregon St at Adeline St
210165          8th St at Ringold St
261132          3rd St at Townsend St
690788          Market St at 10th St

```

```

          end_station_latitude  end_station_longitude  bike_id  user_type \
89574          37.792251          -122.397086          2832  Customer
192814          37.857567          -122.267558          2997  Subscriber
210165          37.774520          -122.409449          5413  Subscriber
261132          37.778742          -122.392741          1672  Subscriber
690788          37.776619          -122.417385          2941  Subscriber

```

```

          bike_share_for_all_trip  member_birth_year  member_gender
89574          No          NaN          NaN
192814          No          1994.0          Female
210165          No          1987.0          Male
261132          No          1977.0          Female
690788          No          NaN          NaN

```

```
In [3]: ford_data.shape
```

```
Out[3]: (870904, 16)
```

```
In [4]: ford_data.describe()
```

```

Out[4]:          duration_sec  start_station_id  start_station_latitude \
count  870904.000000          870174.000000          870904.000000
mean      776.665073          139.027635          37.771409
std      1904.675372          113.704274          0.105744
min        61.000000           3.000000          0.000000

```

25%	337.000000	44.000000	37.770407
50%	537.000000	102.000000	37.780787
75%	840.000000	238.000000	37.797280
max	86114.000000	420.000000	37.880222

	start_station_longitude	end_station_id	end_station_latitude \
count	870904.000000	870174.000000	870904.000000
mean	-122.354787	137.507940	37.770832
std	0.174646	113.653225	0.197522
min	-122.453704	3.000000	0.000000
25%	-122.413004	44.000000	37.770407
50%	-122.398438	100.000000	37.781010
75%	-122.291376	233.000000	37.797320
max	0.000000	420.000000	37.880222

	end_station_longitude	bike_id	member_birth_year
count	870904.000000	870904.000000	175147.000000
mean	-122.352015	4284.292828	1984.806437
std	0.568010	1874.995593	10.116689
min	-122.453704	11.000000	1878.000000
25%	-122.411738	2916.000000	1980.000000
50%	-122.398285	4882.000000	1987.000000
75%	-122.291415	5568.000000	1992.000000
max	0.000000	7108.000000	2001.000000

```
In [5]: ford_data.isnull().sum()
```

```
Out[5]: duration_sec          0
start_time                    0
end_time                      0
start_station_id              730
start_station_name            730
start_station_latitude        0
start_station_longitude       0
end_station_id                730
end_station_name              730
end_station_latitude          0
end_station_longitude         0
bike_id                       0
user_type                     0
bike_share_for_all_trip       0
member_birth_year             695757
member_gender                 695757
dtype: int64
```

```
In [28]: #Create a new column to calculate age of Members
ford_data['member_age'] = 2022 - ford_data['member_birth_year']
```

```
In [39]: #create a copy dataframe for exploration
ford_cleandata = ford_data.copy()
```

```
In [43]: #drop columns not needed for analysis
ford_cleandata.drop(['start_station_id', 'end_station_id', 'start_station_latitude', 'start_station_longitude'], axis=1)
ford_cleandata.sample(10)
```

```
Out[43]:
```

	duration_sec		start_time		end_time	\
275834	1230	2019-02-17	20:33:08.892	2019-02-17	20:53:39.250	
681550	180	2019-04-23	12:55:32.878	2019-04-23	12:58:33.358	
159257	326	2019-01-08	09:22:33.134	2019-01-08	09:27:59.996	
457618	507	2019-03-21	18:33:01.426	2019-03-21	18:41:28.660	
480150	517	2019-03-19	17:07:08.885	2019-03-19	17:15:46.766	
286768	1018	2019-02-15	15:52:27.046	2019-02-15	16:09:25.882	
594894	1084	2019-03-06	17:10:53.872	2019-03-06	17:28:58.785	
173691	711	2019-01-04	16:12:39.862	2019-01-04	16:24:31.520	
359053	638	2019-02-04	17:50:49.115	2019-02-04	18:01:27.768	
673479	911	2019-04-24	13:51:19.369	2019-04-24	14:06:31.311	

	start_station_name	\
275834	Market St at Dolores St	
681550	Washington St at Kearny St	
159257	Folsom St at 3rd St	
457618	Mission Playground	
480150	4th St at Harrison St	
286768	Jackson Playground	
594894	West Oakland BART Station	
173691	The Embarcadero at Steuart St	
359053	2nd St at Folsom St	
673479	Civic Center/UN Plaza BART Station (Market St ...	

	end_station_name	bike_id	\
275834	2nd St at Townsend St	5226	
681550	Clay St at Battery St	1168	
159257	Mechanics Monument Plaza (Market St at Bush St)	4733	
457618	S Van Ness Ave at Market St	6443	
480150	Mississippi St at 17th St	5936	
286768	San Francisco Ferry Building (Harry Bridges Pl...	3223	
594894	West Oakland BART Station	5985	
173691	Berry St at 4th St	5029	
359053	San Francisco Caltrain Station 2 (Townsend St...	3299	
673479	19th St at Florida St	1199	

	user_type	bike_share_for_all_trip	member_gender	member_age
275834	Subscriber	No	Male	47.0
681550	Subscriber	No	NaN	NaN
159257	Subscriber	No	NaN	NaN
457618	Subscriber	No	NaN	NaN
480150	Subscriber	No	NaN	NaN
286768	Subscriber	No	Male	51.0
594894	Subscriber	No	NaN	NaN

173691	Subscriber	No	NaN	NaN
359053	Subscriber	No	Male	30.0
673479	Customer	No	NaN	NaN

```
In [7]: #check for missing data
        ford_cleandata.isnull().sum()
```

```
Out[7]: duration_sec          0
        start_time           0
        end_time             0
        start_station_name    730
        end_station_name      730
        bike_id              0
        user_type            0
        bike_share_for_all_trip 0
        member_gender         695757
        member_age           695757
        dtype: int64
```

```
In [40]: #Drop null rows in start_station_name and end_station_name
         ford_cleandata.dropna(subset=['start_station_name'], inplace=True)
         ford_cleandata.dropna(subset=['end_station_name'], inplace=True)
```

```
In [41]: #change data types for start time, end time
         ford_cleandata['start_time'] = pd.to_datetime(ford_cleandata['start_time'])
         ford_cleandata['end_time'] = pd.to_datetime(ford_cleandata['end_time'])
```

```
In [44]: ford_cleandata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 870174 entries, 0 to 870903
Data columns (total 10 columns):
duration_sec          870174 non-null int64
start_time           870174 non-null datetime64[ns]
end_time             870174 non-null datetime64[ns]
start_station_name    870174 non-null object
end_station_name      870174 non-null object
bike_id              870174 non-null int64
user_type            870174 non-null object
bike_share_for_all_trip 870174 non-null object
member_gender         174952 non-null object
member_age           174952 non-null float64
dtypes: datetime64[ns](2), float64(1), int64(2), object(5)
memory usage: 73.0+ MB
```

```
In [34]: ford_cleandata.shape
```

```
Out[34]: (870174, 10)
```

1.3.1 What is the structure of your dataset?

After basic cleaning the dataset contains 870174 rows and 10 columns: With Missing values in member_age and member_gender, this is a very high number and can not be dropped. 7 columns had been dropped as i would not be needing them for analysis - start_station_id, end_station_id, start_station_latitude, start_station_longitude, end_station_latitude, end_station_longitude, member_birth_year

1.3.2 What is/are the main feature(s) of interest in your dataset?

1.Trip duration (time) 2.Distribution of User Type 3.Monthly bike ride trend 4.Does Age have correlation with how long a rider rides

1.3.3 What features in the dataset do you think will help support your investigation into your feature(s) of interest?


The duration_sec column will provide information we need to calculate the average trip time. The user-type column will be used to find the distribution. Also, we can extract exact month from start_time and end time the duration_sec column will define the average trip time for further analysis. The age of riders has been calculated from the birth_year column of riders.

1.4 Univariate Exploration

1.4.1 Distribution of User Type

In [9]: *#Create a countplot that shows relationship of usertypes*


```
base_color = sb.color_palette()[1]
plot = sb.countplot(data=ford_cleandata, x='user_type', color=base_color)
plt.xlabel('')
plt.ylabel('Number of Bike Trips')
plt.title('Number of bike trip by Subscriber and customer');
#show the plot
plt.show()
```



This shows that the subscribers tend to take more rides than customers of the Ford go bike system

In [13]: *#Create a countplot that shows relationship of riders by gender*

```
base_color = sb.color_palette()[1]
plot = sb.countplot(data=ford_cleandata, x='member_gender', color=base_color)
plt.xlabel('')
plt.ylabel('Number of Bike Trips')
plt.title('Distribution by Gender');
#show the plot
plt.show()
```



A large number of riders are male, while a portion are female, a less portion may have reasons for not revealing their gender.

1.4.2 Age distribution

```
In [14]: #the unique counts of age of riders
         ford_cleandata.member_age.value_counts()
```


```
Out[14]: 34.0    10214
         29.0     9323
         33.0     8967
         32.0     8640
         31.0     8484
         30.0     8245
         35.0     8010
         36.0     7953
         28.0     7654
         27.0     7420
         37.0     7023
         38.0     6557
         39.0     5953
         42.0     5011
         40.0     4987
         26.0     4637
         41.0     4344
         43.0     3756
```


25.0	3476
24.0	3208
44.0	2830
45.0	2706
48.0	2633
23.0	2504
47.0	2503
46.0	2435
49.0	2080
54.0	1927
51.0	1924
50.0	1909
...	
68.0	301
70.0	189
71.0	180
72.0	178
69.0	158
75.0	135
67.0	134
77.0	105
73.0	99
91.0	89
122.0	53
74.0	51
21.0	34
79.0	30
80.0	21
89.0	20
76.0	19
120.0	11
83.0	11
81.0	9
121.0	6
84.0	3
102.0	3
88.0	2
78.0	2
92.0	1
112.0	1
94.0	1
95.0	1
144.0	1

Name: member_age, Length: 75, dtype: int64

```
In [15]: #a graph to show age distribution
plt.figure(figsize=(14,8))
sb.boxplot(x='member_age', data=ford_cleandata, palette='Blues')
```

```
plt.title("Riders age distribution ", fontsize=15, y=1.07)
plt.xlabel("Age", fontsize=15);
```




Outliers are shown so we will drop rows with abnormal age, from 100 and above

```
In [35]: #Drop rows with abnormal age, 100 and above
ford_cleandata =ford_cleandata[ford_cleandata['member_age'] <= 100]
```

```
In [36]: #recheck with a graph to show age distribution
plt.figure(figsize=(14,8))
sb.boxplot(x='member_age', data=ford_cleandata, palette='Blues')
plt.title("Riders age distribution ", fontsize=15, y=1.07)
plt.xlabel("Age", fontsize=15);
```

Riders age distribution



```
In [38]: ford_cleandata.sample(10)
```

```
Out[38]:
```

	duration_sec	start_time	end_time	\
243746	340	2019-02-21 17:53:10.277	2019-02-21 17:58:51.047	
302167	751	2019-02-12 18:13:22.191	2019-02-12 18:25:53.281	
262844	1035	2019-02-19 17:50:57.700	2019-02-19 18:08:13.495	
359828	587	2019-02-04 16:51:49.227	2019-02-04 17:01:36.511	
374105	783	2019-02-01 08:34:39.943	2019-02-01 08:47:43.895	
196312	535	2019-02-28 16:13:34.907	2019-02-28 16:22:30.448	
309388	557	2019-02-11 22:43:10.139	2019-02-11 22:52:27.281	
227920	493	2019-02-23 17:36:45.529	2019-02-23 17:44:58.742	
268548	750	2019-02-19 08:14:14.301	2019-02-19 08:26:44.606	
311213	179	2019-02-11 18:13:27.178	2019-02-11 18:16:26.643	

	start_station_name	\
243746	Downtown Berkeley BART	
302167	Salesforce Transit Center (Natoma St at 2nd St)	
262844	Beale St at Harrison St	
359828	Folsom St at 3rd St	
374105	Howard St at Beale St	
196312	Howard St at 2nd St	
309388	Jones St at Post St	
227920	MLK Jr Way at 14th St	
268548	San Francisco Caltrain (Townsend St at 4th St)	

	end_station_name	bike_id	user_type	\
243746	Bancroft Way at College Ave	5785	Subscriber	
302167	Berry St at King St	5014	Subscriber	
262844	23rd St at Tennessee St	5016	Subscriber	
359828	3rd St at Townsend St	5150	Subscriber	
374105	8th St at Ringold St	987	Subscriber	
196312	San Francisco Caltrain (Townsend St at 4th St)	1832	Subscriber	
309388	16th St Mission BART Station 2	2497	Subscriber	
227920	MLK Jr Way at 14th St	5620	Subscriber	
268548	Howard St at Beale St	2714	Subscriber	
311213	Union St at 10th St	275	Subscriber	

	bike_share_for_all_trip	member_gender	member_age
243746	No	Male	35.0
302167	No	Male	31.0
262844	No	Female	42.0
359828	No	Male	29.0
374105	No	Male	35.0
196312	No	Male	41.0
309388	No	Male	35.0
227920	No	Male	54.0
268548	No	Male	46.0
311213	No	Male	36.0

Most of the riders are ranged from ages 30 - 45 years old which makes sense for long distance races


1.4.3 Monthly bike ride trend

```
In [45]: ford_cleandata['month'] = ford_cleandata.start_time.dt.month
```

```
In [46]: ford_cleandata['month'].value_counts()
```

```
Out[46]: 3    256078
         4    239047
         1    191834
         2    183215
         Name: month, dtype: int64
```


```
In [47]: sb.countplot(data=ford_cleandata, x='month', color=sb.color_palette()[1])
plt.title('Number of Rides per month');
plt.xticks([0, 1, 2, 3], ['Jan', 'Feb', 'March', 'April']);
```



1.4.4 Trip duration (time)

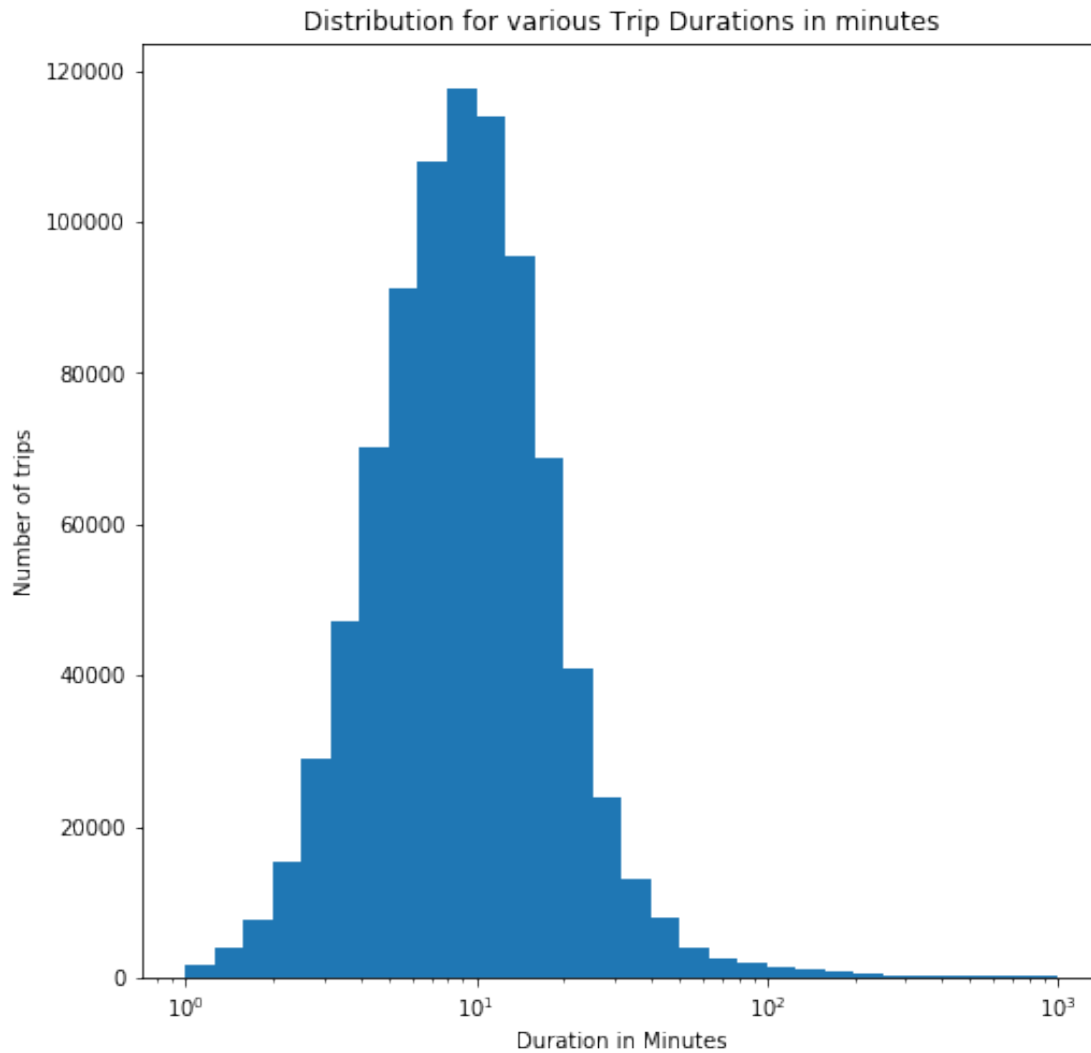
```
In [48]: # Plot a histogram representing time spent on each race in seconds
        binsize = 50
        bins = np.arange(50, ford_cleandata['duration_sec'].max()+binsize, binsize)

        plt.figure(figsize=[8, 5])
        plt.hist(data = ford_cleandata, x = 'duration_sec', bins=bins)
        plt.title('Distribution for various Trip Durations in seconds')
        plt.xlabel('Duration in seconds')
        plt.ylabel('Number of trips')
        plt.xlim([0, 3000]);
```



```
In [49]: #to get the average duration of a races in minutes, we will convert the seconds column
ford_cleandata['race_mins'] = ford_cleandata['duration_sec'] / 60
```

```
plt.figure(figsize=[8,8])
bins = 10*np.arange(0 , 3+0.1 , 0.1)
ticks = [ 0.1 , 0.3 , 1 , 3, 10, 30, 100, 300]
labels = ['{}'.format(v) for v in ticks]
plt.hist(ford_cleandata['race_mins'],bins=bins)
plt.xticks(ticks,labels)
plt.xscale('log')
plt.title('Distribution for various Trip Durations in minutes')
plt.xlabel('Duration in Minutes')
plt.ylabel('Number of trips')
plt.show()
```



1.4.5 Bike Share for all rides


In [33]: *# By plotting a pie graph, we can tell the distribution of bikes.*

```
labels = list(ford_cleandata.bike_share_for_all_trip.unique())
sizes = ford_cleandata.bike_share_for_all_trip.value_counts()
explode = (0.16, 0)

fig1, ax1 = plt.subplots()
ax1.pie(sizes, labels=labels, autopct='%1.1f%%',
        shadow=True, startangle=90, explode = explode)

ax1.axis('equal')
plt.tight_layout()
```

```
plt.title('Bike Share for all trips')
plt.show()
```



A very low proportion of bike riders shared bikes compared to those who weren't.

1.4.6 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?

Age: most of riders age falls between 30 and 45 years old.

Gender: A large number of riders are male, while a portion are female, a less portion may have reasons for not revealing their gender. it needs more investigation because a lot of riders did not indicate their gender, and this would be difficult to determine.

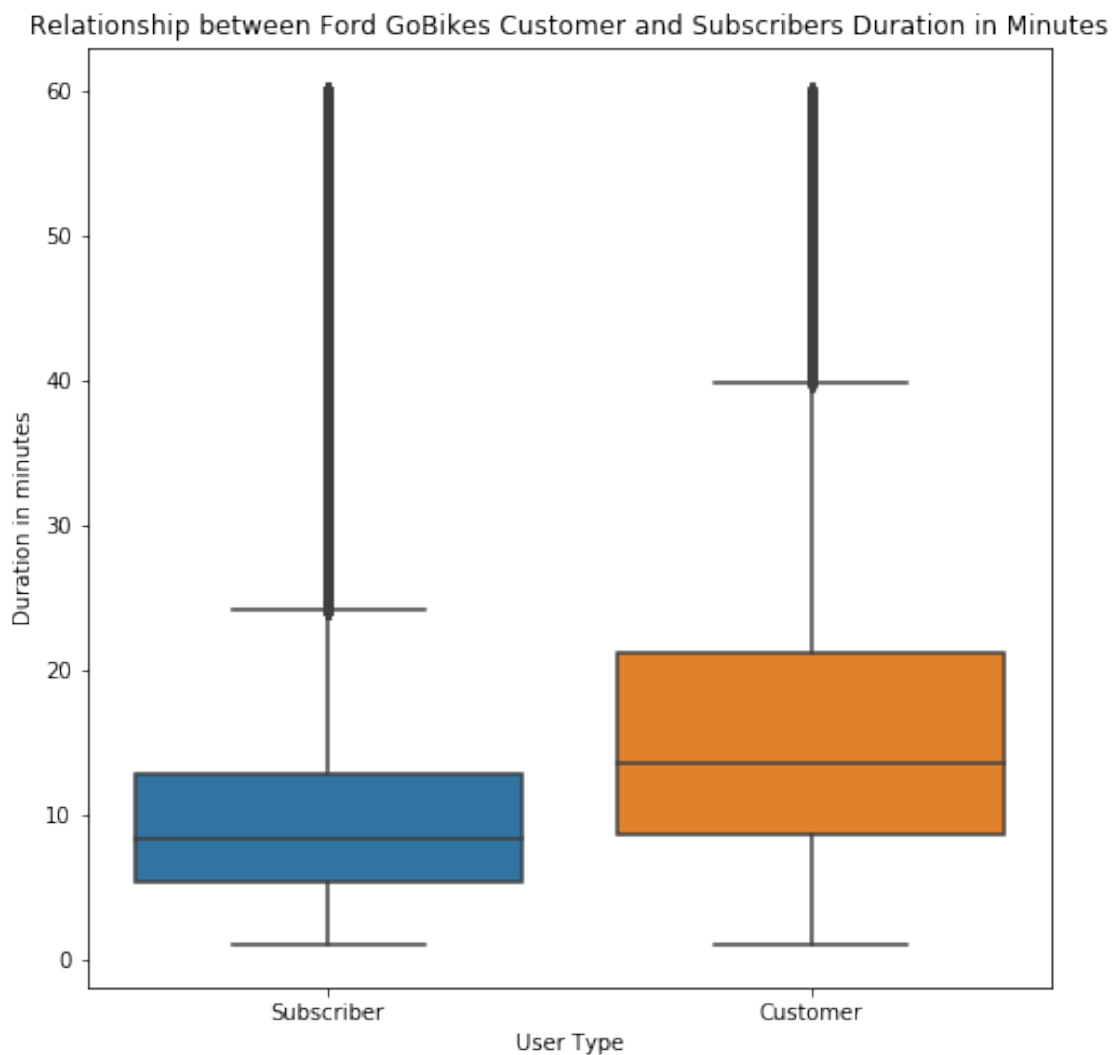
Subscribe: the number of trips in subscribers is more than the number in customers this may be because of pricing and population.

1.4.7 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?

The new column age had a lot of outliers, that has been dealt with by dropping rows with high numbers. New column to calculate the month in which a race was done.

1.5 Bivariate Exploration

```
In [50]: #A boxplot will show the relationship between usertypes and duration of bike trips
plt.figure(figsize = [8, 8])
base_color = sb.color_palette()[3]
sb.boxplot(data = ford_cleandata.query('race_mins <= 60'), x = 'user_type', y = 'race_mins')
plt.title('Relationship between Ford GoBikes Customer and Subscribers Duration in Minutes')
plt.xlabel('User Type')
plt.ylabel('Duration in minutes')
plt.show()
```




Customers had longer bike trips than subscribers


```
In [35]: #use a countplot to show the relationship between the distribution of bike share and the month
sb.countplot(data=ford_cleandata, x='month', hue='bike_share_for_all_trip')
plt.xticks([0, 1, 2, 3], ['Jan', 'Feb', 'March', 'April']);
```

```
plt.xlabel('Months')
plt.ylabel('')
plt.title("Distribution of Bike Share per Month")
```


Out[35]: Text(0.5,1,'Distribution of Bike Share per Month')



```
In [39]: #a countplot to show the relationship between user types and the month of each trip
g = sb.countplot(data=ford_cleandata, x='month', hue='user_type')
g.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.xticks([0, 1, 2, 3], ['Jan', 'Feb', 'March', 'April']);
#plt.xlabel('Months')
plt.ylabel('Number of Users')
plt.title("Distribution of Users' Type per Month");
```



```
In [40]: # Makes the figure enlarged for better visualization
plt.figure(figsize = [10,10])
plt.scatter(data = ford_cleandata , x = 'month' , y = 'race_mins')
plt.xticks([1, 2, 3, 4 ], ['Jan', 'Feb', 'March', 'April']);
plt.title('The Month-wise Trip Duration Trend')
plt.xlabel(' ')
plt.ylabel('Trip Duration in Minutes');
```



1.5.1 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?


Younger riders between the age 20 and 45 tend to take more rides than older people. A high percentage of riders are customers compared to subscribers. Female users average bike trip duration is slightly more than male users. Subscribed users numbers are way more greater than customers despite of the gender.

1.5.2 Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

Trip duration in female users and customer users is longer despite of their low count in the dataset A lot of missing data for gender and age, so we cannot base any conclusion on those two observations

1.6 Multivariate Exploration


```
In [51]: # Find out if Gender affects a trip duration during thr first quater of the year using
graph = sb.FacetGrid(data = ford_cleandata, hue = 'member_gender')
graph.map(plt.scatter, 'month', 'race_mins', alpha = 1/3).add_legend()
plt.xticks([1, 2, 3, 4 ], ['Jan', 'Feb', 'March', 'April'])
plt.title('Trip Duration(min) Against Month and Gender')
plt.xlabel('Month')
plt.ylabel('Trip Duration(min)');
```



Only Feburay bike riders gave gender details so no valid conclusion can be drawn from this

1.6.1 Does duration of trip affects the user types and the month


```
In [42]: #finding how duration affects the user types and the month
sb.pointplot(data = ford_cleandata, x = 'month', y = 'duration_sec', hue = 'user_type',
plt.title('Trip duartion across user type in different months')
plt.xticks([0, 1, 2, 3], ['Jan', 'Feb', 'March', 'April']);
plt.ylabel('Mean Trip Dduration (sec)')
plt.xlabel('Months')
plt.show();
```



The Customers had longer trips and this is spread across the four months

1.6.2 Does age affect how long a biker rode, and whats the relationship with user type also

```
In [52]: plt.figure(figsize=(15,8))
          ax = sb.pointplot(x='member_age', y='race_mins', hue='user_type',data=ford_cleandata.qu
          plt.title('Duration of race by User Type And Age', fontsize=22, y=1.015)
          plt.xlabel('Age', labelpad=16)
          plt.ylabel('Race Duration in Minutes', labelpad=16)
          leg = ax.legend()
          leg.set_title('Gender',prop={'size':16})
          ax = plt.gca();
```



1.7 Conclusions

The dataset is for the first 4 months in the year 2019, January, February, March and April. A large number of riders are male, while a portion are female, a less portion may have reasons for not revealing their gender. Outliers are shown so we will drop rows with abnormal age, from 100 and above. Most of the riders are ranged from ages 30 - 45 years old which makes sense for long distance races. A lot of the bikes rides are covered in short minutes, the average trips takes around 500 - 550 seconds. The age disparity and the duration of trips shows that younger people takes shorter trips compared to older people. A very low proportion of bike riders shared bikes compared to those who weren't. The category user type shows that a over 80% of riders are subscribers as compared to nearly 20% of customers, this may be linked to the prices of hikes and funding system needed to take on a trip. This has a great effect on the number of trips taken and the duration of it. Subscribers takes longer and more trips than customers.

1.8 Recommendation:

Ford go bike should make conscious effort to make more users subscribers. Marketing strategy should be put in place to get new subscribers or convert customers to subscribers.

```
In [53]: ford_cleandata.to_csv('ford_cleandata1.csv')
```

```
In [ ]:
```